

Meera / Skinstinct Content Pipeline — Automation Brief, Nine Checks, and Components Map

Prepared for: MESA AI and its Application, Cohort C4, Case 1

Note: the "Nine Checks" used here are a general-purpose AI-automation risk framework I constructed for this exercise (reversibility, hallucination risk, judgment-substitution, data sensitivity, access control, silent failure, right-layer-of-problem, feedback drift, scope creep). They are not verified against your course's actual taught framework — swap in your instructor's version if it differs before using this to defend in class.

1. Automation Brief

Pain. Meera captures 2–3 raw notes a week in a personal Telegram channel — roughly 60 fragments over eight months — and none of them became a published post. She also has 40 abandoned Google Drive drafts. The posts that did get published took 90 minutes to 3 hours of what she calls "stall time," not writing time: she opens a doc, writes two lines, decides it isn't good enough, and closes it. A hired writer's technically-correct drafts required more rewriting than writing from scratch, so the constraint isn't research effort or grammar — it's activation energy at the blank page, in her own voice, with technically precise content she can stand behind.

User: Meera Pillai — solo founder, sole author of her brand's voice, and the only person who can validate the technical/formulation claims that make her writing credible. She is both the source of raw material (the notes) and the only valid final judge of what goes out under her name.

Outcome. Three LinkedIn posts a week, sourced from the note habit she already has and won't change, without the process consuming her time. Success means a draft reaches her close enough to "post-ready" that her remaining work is a judgment call (approve / discard / revise), not composition from scratch.

The Journey Today. Note captured in Telegram !' sits unopened, disconnected from anything !' separately, she opens a Drive doc to try to write a post !' writes a line or two !' self-doubt !' closes the doc !' both the note and the draft go cold !' repeat, 60 times over. The bottleneck is not idea supply or drafting skill; it's the transition from raw thought to a first version she's willing to look at again.

2. The Nine Checks

Check 1 — Reversibility

Can a wrong action be undone before it costs something?

Evidence: Two no-code consultants proposed end-to-end tools; she rejected both. The one genuinely irreversible action anywhere in this workflow is a LinkedIn post reaching 6,200 followers under her name — credibility she has built over 18 months, tied directly to a 67% repeat-purchase rate she attributes partly to trust in her communication.

Verdict: fails if the bot can publish. Passes only if "publish" stays a manual, human-executed action outside the bot's authority — the bot's furthest action is handing her an approved draft, not posting it.

Check 2 — Hallucination / Fact-Grounding

Can the AI introduce a claim that isn't independently checkable?

Evidence: Meera's authority is built entirely on exact, checkable numbers — pH thresholds, CoA data, percentages (linkedin_post_001, newsletter_003). Section 13 of her voice guide explicitly forbids inventing a stat when real data isn't available.

Verdict: fails for any design where an LLM freely blends note content and news-search results into fluent prose with no traceability back to a source.

Check 3 — Judgment-Substitution

Is the AI being asked to make a call that's actually hers to make?

Evidence: the requested "Gatekeeper Filter" asks the model to unilaterally decide whether eight months of her own raw thinking is "worth developing" — with no described way for her to see or recover what got filtered out.

Verdict: conditional fail. This is the check that produces the Cut, below.

Check 4 — Data Sensitivity

Does note content leave her control in ways she hasn't agreed to?

Evidence: her notes contain unpublished business data — return-rate percentages, an admittedly-

unsolved Vitamin-C formulation attempt, revenue figures — that would pass through at least two third-party APIs (the LLM provider and the news/search provider).

Verdict: conditional pass, provided only the minimal search query (not the raw note) goes to the news API, and the LLM provider is one with a no-training-on-API-data policy.

Check 5 — Access Control

Who can actually trigger this bot?

Evidence: this is her personal Telegram channel; nothing in the case suggests any other user should ever interact with it.

Verdict: fails by default. A bot token alone doesn't restrict senders — it must hard-check the sender's Telegram user ID against an allow-list of exactly one.

Check 6 — Silent Failure

What happens when a dependency breaks?

Evidence: she already lost 60 notes to inaction once. A bot that swallows an API error and does nothing recreates the identical failure mode with more moving parts.

Verdict: fails for any handler without explicit error handling that tells her, in Telegram, that something broke and what state her note is in.

Check 7 — Right Layer of the Problem

Does automating this actually address the stated pain?

Evidence: her stall time is described as psychological/activation-energy ("I know what I want to say. I can't get myself to sit down and shape it"), not a lack of raw material or inability to write technically.

Verdict: conditional. Auto-drafting is the correct layer only if draft quality is high enough that revising feels like editing, not restarting. A mediocre AI draft in a flat imitation of her voice just becomes a new blank page dressed differently.

Check 8 — Feedback Loop / Drift

Does the system get better from her approvals and discards, or stay static?

Evidence: nothing in the requested design has the bot learn from her decisions over time.

Verdict: fails, acceptable as a known v1 limitation — but every approve/discard/feedback-text event should be logged now, even unused, so a later version can learn from it.

Check 9 — Scope Creep ("the whole thing")

Is any part of this doing more than she actually asked for?

Evidence: she has already rejected two "does everything end-to-end" tools.

Verdict: this check re-runs on every future feature request. No auto-posting, auto-scheduling, or cross-platform posting should be added just because it's technically easy — only because she asks for it.

3. The Cut

The one thing she's described wanting that a check stops you from building as originally specified: fully autonomous news/data enrichment woven into finished prose as if it were already verified, with no visible source and no checkpoint before the draft reaches her.

She explicitly wants the draft to "reference something current... so the post doesn't feel like it was written in a vacuum" — a legitimate ask (Check 7). But her voice, and her brand's entire credibility

model, is built on every number being independently checkable (Checks 1 and 2). A search API can return a misdated or misquoted source; an LLM can paraphrase a stat and drop a qualifier or shift a decimal. If that error reaches 6,200 followers under her name, it isn't a typo — it's the exact "label says X, reality is more complicated" gap her whole brand exists to call out, now pointed at her.

The Cut, concretely: the bot still enriches with news — that feature stays — but never silently.

Every externally-sourced fact in a draft carries a visible inline citation and link, is marked unverified, and the Approve button's own confirmation copy reminds her to check cited facts before she actually posts. Approving a draft in the bot is not the same event as fact-checking it, and the tool should never imply otherwise. Nothing about her *reach* (three posts a week, low personal effort) is cut — only the false confidence that an AI-fetched fact needs no human eyes before it goes out as hers.

4. Components Map

Trigger — a new Telegram message arrives from Meera's own Telegram user ID, in her personal bot chat.

Input — the raw message text: either a new note, a feedback on post draft: {text} revision instruction, or an Approve/Discard button callback.

Context —

- meera_voice.txt: persona, tone rules, structural blueprints, phrase bank, the Section-14 pre-published list
- Draft state: is there a draft currently pending on this note? (needed to route feedback vs. new-note handling and to know what to revise)
- New-note handling: the single checkable arg is extracted from the note (not the whole note) is what gets searched.
- Processing path:
 1. Route: message check / send feedback must match the tone allowed; user, otherwise ignore!
 2. New-note path: gatekeeper classification against her core beliefs (voice-file Section 2) — viable / not viable; plus the specific apply if viable, store status back log, stop
 3. Generate the draft: voice file to be ignored search vs search explicit anti-hallucination instruction!
 4. Draft revision path: add the pending draft, as source note, and its citations; regenerate using her feedback text as the edit instruction, same voice constraints.
 5. AI — two distinct calls against the same LLM provider, both grounded in the same versioned voice-file system prompt, so they stay auditable together.
- Draft revision call: final post-generation of edit, viable or viable, instruct with a new line reason number.
- Output:
 - Draft message (exact copy) for rejected notes
 - keyboard [Approve] [Discard], plus a visible "Source information" footer for any externally fetched facts before reposting, plus a DB status update
 - Discard: archive copy and linked DB status update.